

Retail_Sales_Analysis

August 25, 2021

1 Retail Sales Profit Analysis

```
[1]: import pandas as pd
```

```
[3]: df=pd.read_csv('Retail_SalesMarketing_ProfitCost.csv')
df.shape
```

```
[3]: (84672, 16)
```

```
[4]: df.head()
```

```
[4]:   Year      Product line Product type      Product \
0  2015  Camping Equipment  Cooking Gear  TrailChef Water Bag
1  2015  Camping Equipment  Cooking Gear  TrailChef Water Bag
2  2015  Camping Equipment  Cooking Gear  TrailChef Water Bag
3  2015  Camping Equipment  Cooking Gear  TrailChef Water Bag
4  2015  Camping Equipment  Cooking Gear  TrailChef Water Bag
```

```
   Order method type Retailer country  Revenue  Planned revenue \
0      Telephone      United States  315,044          437,477
1      Telephone           Canada    13,445          14,313
2      Telephone           Mexico         NaN             NaN
3      Telephone           Brazil         NaN             NaN
4      Telephone           Japan   181,120          235,237
```

```
   Product cost  Quantity  Unit cost  Unit price  Gross profit \
0      158,372    66,385         3.0         7      156,673
1         6,299     2,172         3.0         7         7,146
2           NaN         NaN         NaN         NaN         NaN
3           NaN         NaN         NaN         NaN         NaN
4      89,413    35,696         3.0         7      91,707
```

```
   Unit sale price  Unnamed: 14  Unnamed: 15
0                5           NaN           NaN
1                6           NaN           NaN
2               NaN           NaN           NaN
3               NaN           NaN           NaN
4                5           NaN           NaN
```

```
[5]: df.drop(columns=['Unnamed: 14', 'Unnamed: 15'], inplace=True)
```

```
[6]: df.columns
```

```
[6]: Index(['Year', 'Product line', 'Product type', 'Product', 'Order method type',  
         'Retailer country', 'Revenue', 'Planned revenue', 'Product cost',  
         'Quantity', 'Unit cost', 'Unit price', 'Gross profit',  
         'Unit sale price'],  
        dtype='object')
```

1.0.1 Problem: Examine whether or not the retail store outdoor products are growing, maturing or declining. Compare your prospects for future sales with past performance by analyzing the industry average.

1.1 Data Cleaning

1.1.1 Lets clean up the data !!

Let's remove the spaces from the column names to remove ambiguity and better efficiency using strip function

```
[8]: df.columns = [cl.strip() for cl in df.columns.tolist()]  
  
df.columns= df.columns.str.replace(' ','_')  
df.columns.tolist()
```

```
[8]: ['Year',  
      'Product_line',  
      'Product_type',  
      'Product',  
      'Order_method_type',  
      'Retailer_country',  
      'Revenue',  
      'Planned_revenue',  
      'Product_cost',  
      'Quantity',  
      'Unit_cost',  
      'Unit_price',  
      'Gross_profit',  
      'Unit_sale_price']
```

```
[9]: df.dtypes
```

```
[9]: Year                int64  
     Product_line       object  
     Product_type       object  
     Product            object  
     Order_method_type  object  
     Retailer_country   object
```

```

Revenue                object
Planned_revenue        object
Product_cost           object
Quantity              object
Unit_cost              float64
Unit_price             object
Gross_profit           object
Unit_sale_price        object
dtype: object

```

Since we have , in between 315,044 etc, it takes it as an object...so let us remove the commas

```
[10]: df.replace(',', '', regex=True, inplace=True)
```

```
[11]: df.head(3)
```

```
[11]:
```

	Year	Product_line	Product_type	Product	\
0	2015	Camping Equipment	Cooking Gear	TrailChef Water Bag	
1	2015	Camping Equipment	Cooking Gear	TrailChef Water Bag	
2	2015	Camping Equipment	Cooking Gear	TrailChef Water Bag	

	Order_method_type	Retailer_country	Revenue	Planned_revenue	Product_cost	\
0	Telephone	United States	315044	437477	158372	
1	Telephone	Canada	13445	14313	6299	
2	Telephone	Mexico	NaN	NaN	NaN	

	Quantity	Unit_cost	Unit_price	Gross_profit	Unit_sale_price
0	66385	3.0	7	156673	5
1	2172	3.0	7	7146	6
2	NaN	NaN	NaN	NaN	NaN

In Pandas missing data is represented by two value:

- None: None is a Python singleton object that is often used for missing data in Python code.
- NaN : NaN (an acronym for Not a Number), is a special floating-point value recognized by all systems that use the standard IEEE floating-point representation (NaN is numpy type)

We find using isnull() and not notnull()

```
[13]: null_series=df.isnull().sum()
# type(null_df)
null_series
```

```
[13]: Year                0
Product_line            0
Product_type            0
Product                0
Order_method_type       0
Retailer_country        0
```

```

Revenue          59929
Planned_revenue  59929
Product_cost     59929
Quantity         59929
Unit_cost        59929
Unit_price       59929
Gross_profit     59929
Unit_sale_price  59929
dtype: int64

```

These are the null/missing values in the dataset

```

[14]: notnull_series=df.notnull().sum()
      # type(notnull_df)
      notnull_series

```

```

[14]: Year          84672
      Product_line  84672
      Product_type  84672
      Product       84672
      Order_method_type 84672
      Retailer_country 84672
      Revenue       24743
      Planned_revenue 24743
      Product_cost  24743
      Quantity      24743
      Unit_cost     24743
      Unit_price    24743
      Gross_profit  24743
      Unit_sale_price 24743
      dtype: int64

```

These are the values that are Not-Null. Let us combine the null and not null to find percentage of missing values in each column

We are concatenating two series null & not null ,into a dataframe vertically (columns) by setting axis =1 and naming them . Missing_df.columns[list] we are assigning the names

```

[15]: missing_df=pd.concat([null_series,notnull_series], axis=1).reset_index()
      missing_df.columns=['Column_name','Null_count','Not_Null_Count']

```

```

[16]: missing_df['Total_count']= missing_df['Null_count'] +_
      ↪missing_df['Not_Null_Count']

      missing_df['Missing_values_percent']= (missing_df['Null_count']/_
      ↪missing_df['Total_count']) *100
      missing_df
      # Null_count/ (Null_count+Not_Null_Count)* 100

```

```
[16]:
```

	Column_name	Null_count	Not_Null_Count	Total_count	\
0	Year	0	84672	84672	
1	Product_line	0	84672	84672	
2	Product_type	0	84672	84672	
3	Product	0	84672	84672	
4	Order_method_type	0	84672	84672	
5	Retailer_country	0	84672	84672	
6	Revenue	59929	24743	84672	
7	Planned_revenue	59929	24743	84672	
8	Product_cost	59929	24743	84672	
9	Quantity	59929	24743	84672	
10	Unit_cost	59929	24743	84672	
11	Unit_price	59929	24743	84672	
12	Gross_profit	59929	24743	84672	
13	Unit_sale_price	59929	24743	84672	

	Missing_values_percent
0	0.000000
1	0.000000
2	0.000000
3	0.000000
4	0.000000
5	0.000000
6	70.777825
7	70.777825
8	70.777825
9	70.777825
10	70.777825
11	70.777825
12	70.777825
13	70.777825

We can find that we have a 70% missing data from column 6 to column 13...which is a lot. We have to eventually delete the data, but let's see which subgroup has maximum no. of missing values

```
[17]: df.groupby(['Product_line'])['Revenue', 'Year'].count()
```

```
<ipython-input-17-753fef286776>:1: FutureWarning: Indexing with multiple keys
(implicitly converted to a tuple of keys) will be deprecated, use a list
instead.
```

```
df.groupby(['Product_line'])['Revenue', 'Year'].count()
```

```
[17]:
```

	Revenue	Year
Product_line		
Camping Equipment	8375	24108
Golf Equipment	2763	8820
Mountaineering Equipment	2947	12348
Outdoor Protection	2944	8820

Personal Accessories 7714 30576

We can see that the null count is distributed across all the product lines, (we narrowed down to check if a particular sub-granular group like product line had lots of missing data wrt only one). We can proceed with deleting them.

There was a **70%** data loss happening across all metrics which cannot be manipulated by filling values using **Fillna**.

```
[18]: df=df.dropna()
```

```
[19]: df.head(3)
```

```
[19]:   Year      Product_line Product_type      Product \
0  2015  Camping Equipment  Cooking Gear  TrailChef Water Bag
1  2015  Camping Equipment  Cooking Gear  TrailChef Water Bag
4  2015  Camping Equipment  Cooking Gear  TrailChef Water Bag

      Order_method_type  Retailer_country  Revenue  Planned_revenue  Product_cost \
0          Telephone      United States  315044          437477        158372
1          Telephone           Canada    13445          14313         6299
4          Telephone           Japan    181120          235237        89413

      Quantity  Unit_cost  Unit_price  Gross_profit  Unit_sale_price
0      66385         3.0         7         156673             5
1       2172         3.0         7           7146             6
4      35696         3.0         7          91707             5
```

```
[20]: df.reset_index(drop= True, inplace=True)

      # Drop= true as we dont want to carry forward that index. Index was 0,1,4
```

```
[21]: df.head(3)
```

```
[21]:   Year      Product_line Product_type      Product \
0  2015  Camping Equipment  Cooking Gear  TrailChef Water Bag
1  2015  Camping Equipment  Cooking Gear  TrailChef Water Bag
2  2015  Camping Equipment  Cooking Gear  TrailChef Water Bag

      Order_method_type  Retailer_country  Revenue  Planned_revenue  Product_cost \
0          Telephone      United States  315044          437477        158372
1          Telephone           Canada    13445          14313         6299
2          Telephone           Japan    181120          235237        89413

      Quantity  Unit_cost  Unit_price  Gross_profit  Unit_sale_price
0      66385         3.0         7         156673             5
1       2172         3.0         7           7146             6
2      35696         3.0         7          91707             5
```

```
[22]: df.shape
```

```
[22]: (24743, 14)
```

```
[23]: df.dtypes
```

```
[23]: Year                int64
      Product_line       object
      Product_type       object
      Product            object
      Order_method_type  object
      Retailer_country   object
      Revenue            object
      Planned_revenue    object
      Product_cost       object
      Quantity           object
      Unit_cost          float64
      Unit_price         object
      Gross_profit       object
      Unit_sale_price    object
      dtype: object
```

Let us strip extra spaces, - and such noises in the columns...

```
[24]: df=df.assign(Revenue= df['Revenue'].str.strip(),
                        Planned_revenue= df['Planned_revenue'].str.strip(),
                        Product_cost= df['Product_cost'].str.strip(),
                        Quantity= df['Quantity'].str.strip(),
                        Unit_price= df['Unit_price'].str.strip(),
                        Gross_profit= df['Gross_profit'].str.strip(),
                        Unit_sale_price= df['Unit_sale_price'].str.strip()
                      )
      # df.iloc[100]
      # df['Revenue'].iloc[100]
```

```
[25]: df=df.assign(Revenue= df['Revenue'].str.replace('-', ''),
                        Planned_revenue= df['Planned_revenue'].str.replace('-', ''),
                        Product_cost= df['Product_cost'].str.replace('-', ''),
                        Quantity= df['Quantity'].str.replace('-', ''),
                        Unit_price= df['Unit_price'].str.replace('-', ''),
                        Gross_profit= df['Gross_profit'].str.replace('-', ''),
                        Unit_sale_price= df['Unit_sale_price'].str.replace('-', '')
                      )
```

Converting to numeric types....

```
[26]: df= df.assign(
      Revenue= pd.to_numeric(df['Revenue']),
```

```

Planned_revenue= pd.to_numeric(df['Planned_revenue']),
Product_cost= pd.to_numeric(df['Product_cost']),
Quantity= pd.to_numeric(df['Quantity']),
Unit_price= pd.to_numeric(df['Unit_price']),
#     Gross_profit= pd.to_numeric(df['Gross_profit']),
Unit_sale_price= pd.to_numeric(df['Unit_sale_price'])
)

```

```
[27]: df.Gross_profit.iloc[301]
```

```
[27]: '(1119)'
```

Since it has paranthesis in place for negative values, lets convert them and change data type

```
[28]: df['Gross_profit']= (df['Gross_profit'].replace( '[\$,)]', '', regex=True )
      .replace( '([,','- ', regex=True ).astype(float))
```

```
[29]: df.dtypes
```

```
[29]: Year                int64
      Product_line        object
      Product_type        object
      Product             object
      Order_method_type    object
      Retailer_country      object
      Revenue             float64
      Planned_revenue      int64
      Product_cost         int64
      Quantity            int64
      Unit_cost            float64
      Unit_price           int64
      Gross_profit         float64
      Unit_sale_price      float64
      dtype: object
```

```
[30]: df.isnull().sum()
```

```
[30]: Year                0
      Product_line        0
      Product_type        0
      Product             0
      Order_method_type    0
      Retailer_country      0
      Revenue             76
      Planned_revenue      0
      Product_cost         0
      Quantity            0
      Unit_cost            0
```



```
Unit_price      0
Gross_profit     0
Unit_sale_price  76
dtype: int64
```

Revenue and Unit sale price still have 76 null values...lets fill them with 0 as they are not missing. They have no revenue and no sale price

```
[31]: df=df.fillna(0)
```

```
[32]: df.isnull().sum()
```

```
[32]: Year      0
Product_line  0
Product_type  0
Product       0
Order_method_type  0
Retailer_country  0
Revenue       0
Planned_revenue  0
Product_cost  0
Quantity      0
Unit_cost     0
Unit_price    0
Gross_profit  0
Unit_sale_price  0
dtype: int64
```

```
[33]: df.columns
```

```
[33]: Index(['Year', 'Product_line', 'Product_type', 'Product', 'Order_method_type',
        'Retailer_country', 'Revenue', 'Planned_revenue', 'Product_cost',
        'Quantity', 'Unit_cost', 'Unit_price', 'Gross_profit',
        'Unit_sale_price'],
        dtype='object')
```

```
[37]: df.Product_type.unique()
```

```
[37]: array(['Cooking Gear', 'Tents', 'Sleeping Bags', 'Packs', 'Lanterns',
        'Watches', 'Eyewear', 'Knives', 'Binoculars', 'Navigation',
        'Insect Repellents', 'Sunscreen', 'First Aid', 'Irons', 'Woods',
        'Putters', 'Golf Accessories', 'Rope', 'Safety',
        'Climbing Accessories', 'Tools'], dtype=object)
```

```
[38]: df.Product.unique()
```

```
[38]: array(['TrailChef Water Bag', 'TrailChef Canteen',
        'TrailChef Kitchen Kit', 'TrailChef Cup', 'TrailChef Cook Set',
```

```

'TrailChef Deluxe Cook Set', 'TrailChef Single Flame',
'TrailChef Double Flame', 'TrailChef Kettle', 'TrailChef Utensils',
'Star Lite', 'Star Dome', 'Star Gazer 2', 'Star Gazer 3',
'Star Gazer 6', 'Star Peg', 'Hibernator Lite', 'Hibernator',
'Hibernator Extreme', 'Hibernator Self - Inflating Mat',
'Hibernator Pad', 'Hibernator Pillow', 'Hibernator Camp Cot',
'Canyon Mule Climber Backpack', 'Canyon Mule Weekender Backpack',
'Canyon Mule Journey Backpack', 'Canyon Mule Extreme Backpack',
'Canyon Mule Cooler', 'Canyon Mule Carryall', 'Firefly Lite',
'Firefly Mapreader', 'Firefly 2', 'Firefly 4', 'Firefly Extreme',
'Firefly Multi-light', 'EverGlow Single', 'EverGlow Double',
'EverGlow Kerosene', 'EverGlow Butane', 'EverGlow Lamp',
'Flicker Lantern', 'Mountain Man Analog', 'Mountain Man Digital',
'Mountain Man Deluxe', 'Mountain Man Combination',
'Mountain Man Extreme', 'Venue', 'Infinity', 'Lux', 'Sam', 'TX',
'Legend', 'Kodiak', 'Polar Sun', 'Polar Ice', 'Polar Sports',
'Polar Wave', 'Polar Extreme', 'Bella', 'Capri', 'Cat Eye',
'Dante', 'Fairway', 'Inferno', 'Maximus', 'Trendi', 'Zone',
'Hawk Eye', 'Single Edge', 'Double Edge', 'Edge Extreme',
'Bear Edge', 'Bear Survival Edge', 'Max Gizmo', 'Pocket Gizmo',
'Seeker 35', 'Seeker 50', 'Seeker Extreme', 'Seeker Mini',
'Opera Vision', 'Ranger Vision', 'Glacier Basic', 'Glacier Deluxe',
'Glacier GPS', 'Glacier GPS Extreme', 'Trail Master',
'Trail Scout', 'Trail Star', 'BugShield Natural',
'BugShield Spray', 'BugShield Lotion Lite', 'BugShield Lotion',
'BugShield Extreme', 'Sun Blocker', 'Sun Shelter Stick',
'Sun Shelter 15', 'Sun Shelter 30', 'Sun Shield',
'Compact Relief Kit', 'Deluxe Family Relief Kit',
'Calamine Relief', 'Aloe Relief', 'Insect Bite Relief',
'Hailstorm Steel Irons', 'Hailstorm Titanium Irons',
'Lady Hailstorm Steel Irons', 'Lady Hailstorm Titanium Irons',
'Hailstorm Titanium Woods Set', 'Hailstorm Steel Woods Set',
'Lady Hailstorm Titanium Woods Set',
'Lady Hailstorm Steel Woods Set', 'Course Pro Putter',
'Blue Steel Putter', 'Blue Steel Max Putter',
'Course Pro Golf and Tee Set', 'Course Pro Umbrella',
'Course Pro Golf Bag', 'Course Pro Gloves', 'Zodiak', 'Retro',
'Astro Pilot', 'Sky Pilot', 'Husky Rope 50', 'Husky Rope 60',
'Husky Rope 100', 'Husky Rope 200', 'Granite Climbing Helmet',
'Husky Harness', 'Husky Harness Extreme', 'Granite Signal Mirror',
'Granite Carabiner', 'Granite Belay', 'Granite Pulley',
'Firefly Climbing Lamp', 'Firefly Charger',
'Firefly Rechargeable Battery', 'Granite Chalk Bag', 'Granite Ice',
'Granite Hammer', 'Granite Shovel', 'Granite Grip', 'Granite Axe',
'Granite Extreme', 'Auto Pilot'], dtype=object)

```

```
[36]: df.Product_line.unique()
```

```
[36]: array(['Camping Equipment', 'Personal Accessories', 'Outdoor Protection',
        'Golf Equipment', 'Mountaineering Equipment'], dtype=object)
```

```
[39]: df.head(3)
```

```
[39]:   Year      Product_line Product_type      Product \
0  2015  Camping Equipment  Cooking Gear  TrailChef Water Bag
1  2015  Camping Equipment  Cooking Gear  TrailChef Water Bag
2  2015  Camping Equipment  Cooking Gear  TrailChef Water Bag

   Order_method_type  Retailer_country  Revenue  Planned_revenue  Product_cost \
0      Telephone      United States  315044.0          437477          158372
1      Telephone          Canada    13445.0           14313           6299
2      Telephone          Japan    181120.0          235237          89413

   Quantity  Unit_cost  Unit_price  Gross_profit  Unit_sale_price
0      66385         3.0          7      156673.0          5.0
1       2172         3.0          7       7146.0          6.0
2      35696         3.0          7      91707.0          5.0
```

1.1.2 The entire data set is cleaned...let us do the exploratory analysis

1.2 Exploratory Analysis

- Summary statistics
- Distribution of Records across product lines
- Distribution of data across countries
- Feature selection Conducting correlation to find which is the best parameter that contri butees to sales profits/ highest revenue

```
[292]: # Summary statistics
```

```
df.describe()
```


```
[292]:   count      Year      Revenue  Planned_revenue  Product_cost \
count  24743.000000  2.474300e+04  2.474300e+04  2.474300e+04
mean    2016.345067  1.894183e+05  1.988176e+05  1.116252e+05
std       1.073106  3.907509e+05  4.025355e+05  2.384156e+05
min     2015.000000  0.000000e+00  1.600000e+01  6.000000e+00
25%     2015.000000  1.857900e+04  1.955700e+04  9.431500e+03
50%     2016.000000  5.986700e+04  6.390700e+04  3.278400e+04
75%     2017.000000  1.901930e+05  2.039955e+05  1.113710e+05
max     2018.000000  1.005429e+07  1.005429e+07  6.756853e+06

   Quantity  Unit_cost  Unit_price  Gross_profit \
count  24743.000000  24743.000000  24743.000000  2.474300e+04
mean    3606.559067    84.946530    156.056541  7.779312e+04
std     8777.721091   131.108962    246.805361  1.581223e+05
```

min	1.000000	1.000000	2.000000	-1.816000e+04
25%	328.000000	11.000000	23.000000	8.333000e+03
50%	1043.000000	37.000000	67.000000	2.579400e+04
75%	3288.000000	80.000000	148.000000	7.825400e+04
max	313628.000000	690.000000	1360.000000	3.521098e+06

	Unit_sale_price
count	24743.000000
mean	147.254900
std	232.045043
min	0.000000
25%	20.000000
50%	63.000000
75%	141.000000
max	1308.000000

```
[51]: import matplotlib.pyplot as plt
fig, axs = plt.subplots(figsize=(8,8))
axs.bar(df['Order_method_type'], df.index)
axs.set_title('Order_method_type')
fig.autofmt_xdate()
```



Revenue over the years

```
[62]: import matplotlib.pyplot as plt
def group_analysis(column_name):
    # Generate the results based on the given column_name
    results= df.groupby(column_name).sum()
    results['Revenue_missed_pct'] =
    →((results['Planned_revenue']-results['Revenue'])/
    →(results['Planned_revenue']))*100.0

    # Then plot the KPI (Key Performance Indicies) for the considered_
    →column_name
    fig, axs = plt.subplots(2, 2, figsize=(15,15))
    axs[0,0].bar(results.index,results['Revenue'])
    axs[0,0].set_title('Revenue')
```

```

fig.autofmt_xdate()

axs[0,1].bar(results.index,results['Planned_revenue'])
axs[0,1].set_title('Planned_revenue')
fig.autofmt_xdate()

axs[1,0].bar(results.index,results['Product_cost'])
axs[1,0].set_title('Product_cost')
fig.autofmt_xdate()

axs[1,1].bar(results.index,results['Gross_profit'])
axs[1,1].set_title('Gross_profit')
fig.autofmt_xdate()

#      Plot for missed revenue and the column name
results.plot.barh(y='Revenue_missed_pct')
plt.ylabel('Revenue_missed_pct')
plt.xlabel=(results.index)

#      ax = results.plot.barh(x='Revenue_missed_pct', y=results.index)
return results

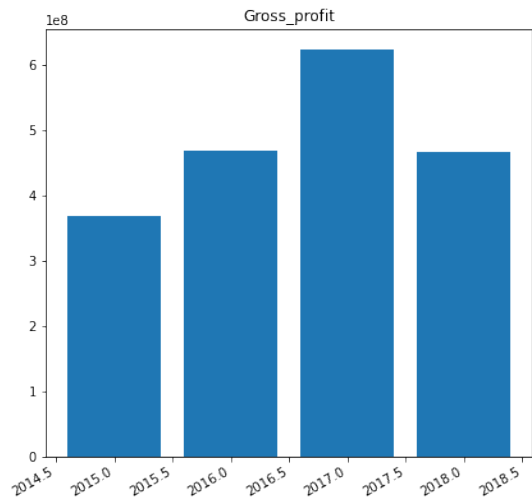
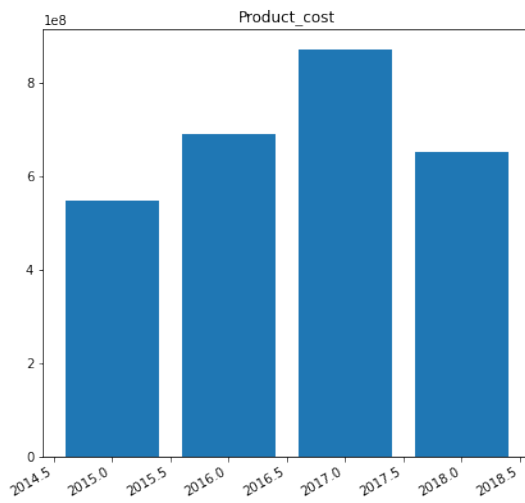
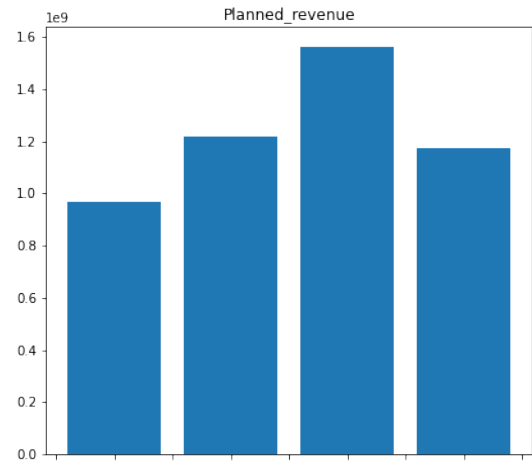

```

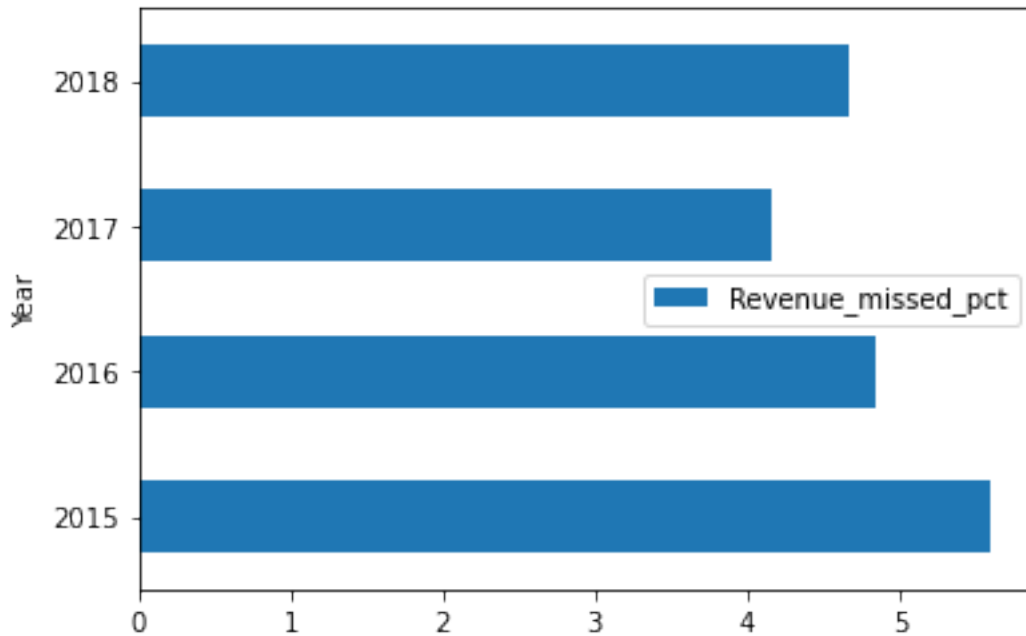
```
[63]: group_analysis('Year')
```

```
[63]:
```

	Revenue	Planned_revenue	Product_cost	Quantity	Unit_cost	\
Year						
2015	9.143529e+08	968475273	546735041	20174730	605714.0	
2016	1.159196e+09	1218177535	690512038	23524685	605797.0	
2017	1.495891e+09	1560738939	872630604	25941790	495691.0	
2018	1.117336e+09	1171950964	652063489	19595886	394630.0	

	Unit_price	Gross_profit	Unit_sale_price	Revenue_missed_pct
Year				
2015	1115154	367617948.0	1042079.0	5.588411
2016	1092246	468683778.0	1031078.0	4.841811
2017	921617	623260635.0	877109.0	4.154940
2018	732290	465272886.0	693262.0	4.660142






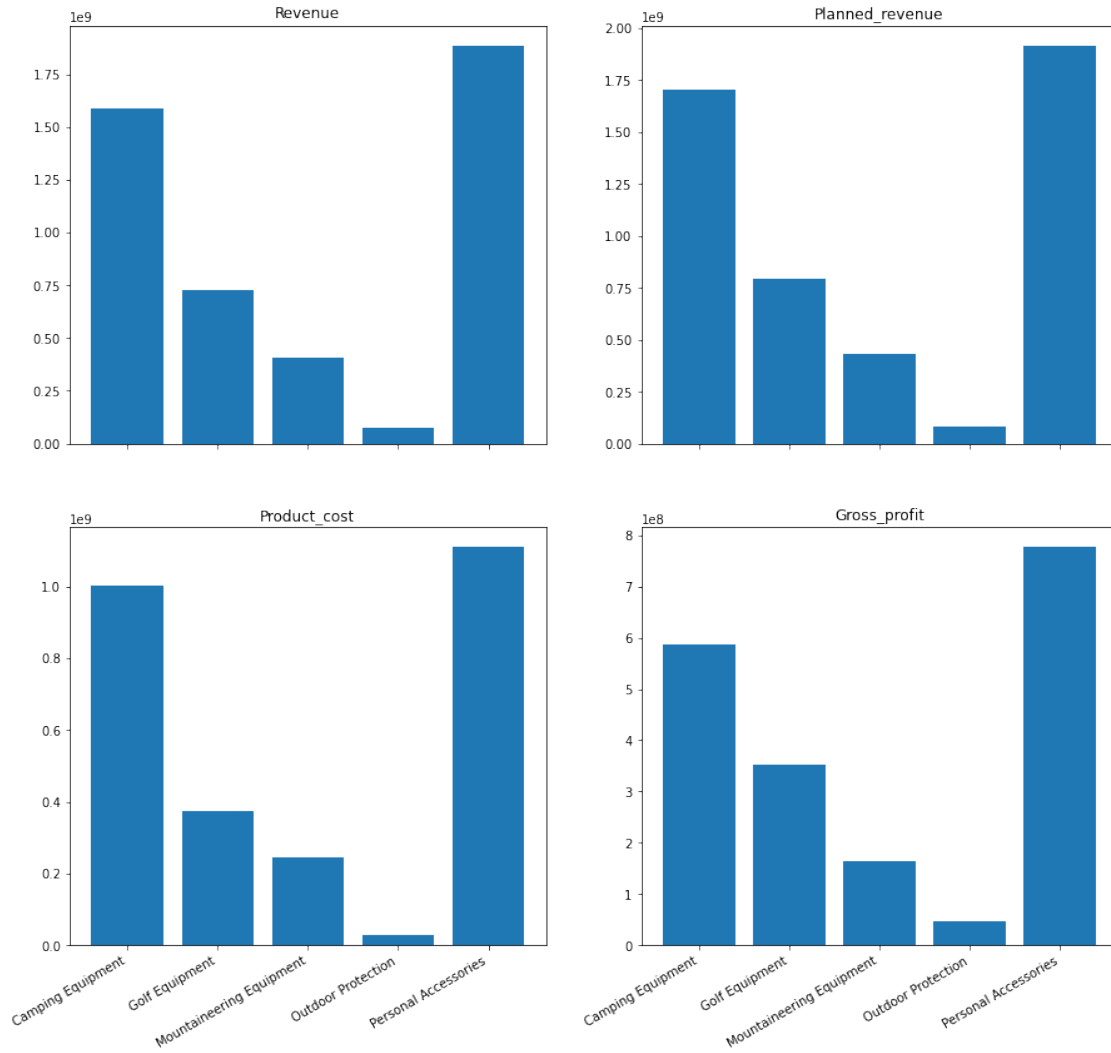
```
[64]: group_analysis('Product_line')
```

```
[64]:
```

	Year	Revenue	Planned_revenue	\
Product_line				
Camping Equipment	16886376	1.589037e+09	1703124790	
Golf Equipment	5571060	7.264114e+08	790261053	
Mountaineering Equipment	5943678	4.096602e+08	430568075	
Outdoor Protection	5935833	7.599434e+07	80005280	
Personal Accessories	15553479	1.885673e+09	1915383513	

	Product_cost	Quantity	Unit_cost	Unit_price	\
Product_line					
Camping Equipment	1002237787	27301149	754776.0	1268118	
Golf Equipment	374217745	5113701	693534.0	1419337	
Mountaineering Equipment	246384239	9900091	185392.0	302861	
Outdoor Protection	30011034	12014445	10472.0	25602	
Personal Accessories	1109090367	34907705	457658.0	845389	

	Gross_profit	Unit_sale_price	Revenue_missed_pct
Product_line			
Camping Equipment	586799155.0	1191877.0	6.698751
Golf Equipment	352193680.0	1322223.0	8.079568
Mountaineering Equipment	163275949.0	287501.0	4.855878
Outdoor Protection	45983327.0	24348.0	5.013339
Personal Accessories	776583136.0	817579.0	1.551127



```
[65]: group_analysis('Retailer_country')
```

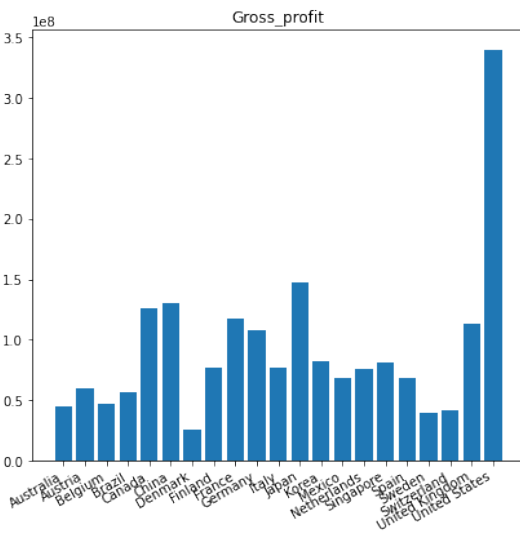
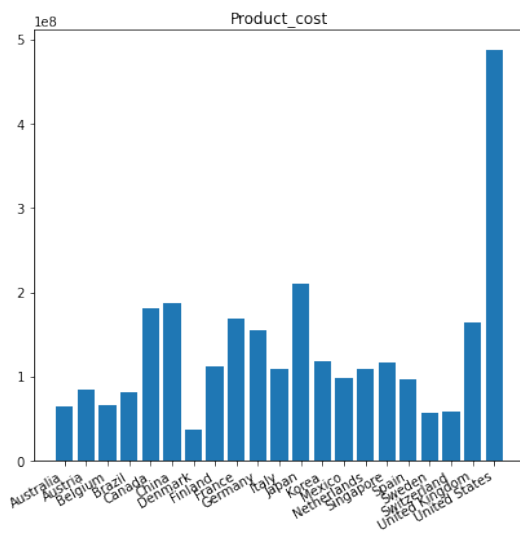

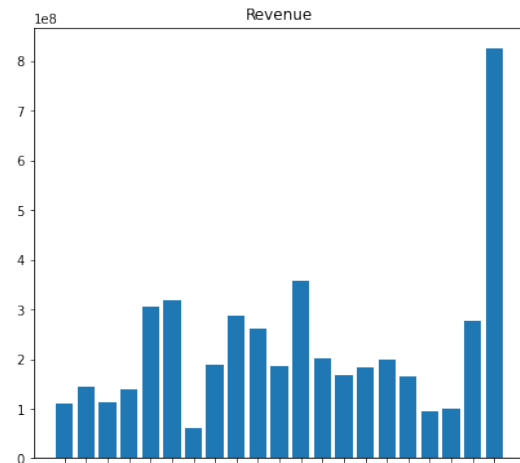
```
[65]:
```

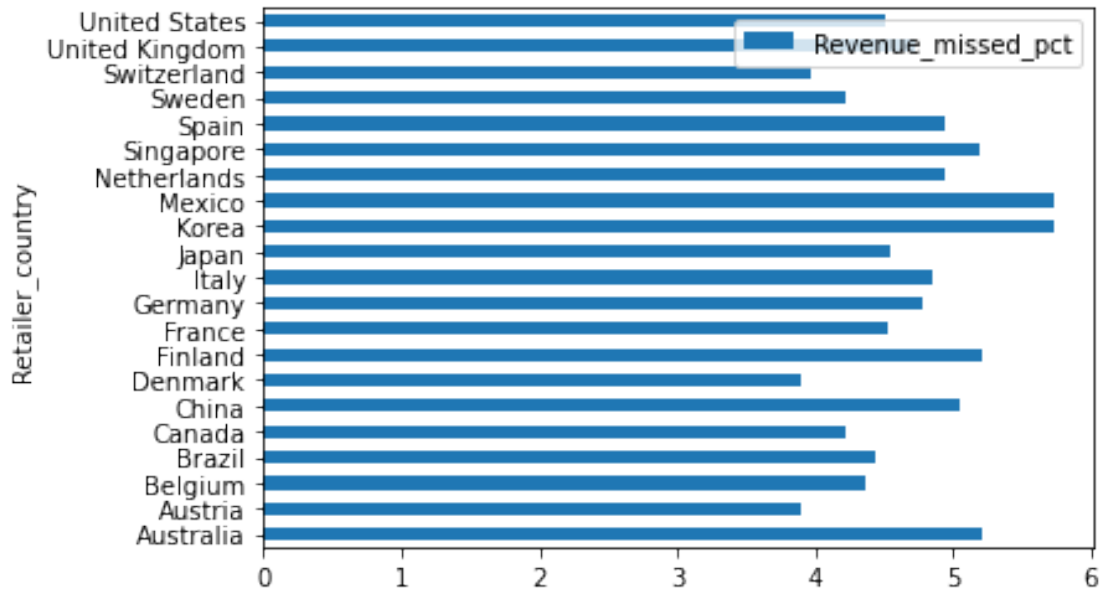
	Year	Revenue	Planned_revenue	Product_cost	\
Retailer_country					
Australia	1857486	109299974.0	115311829	64254527	
Austria	2240195	143709418.0	149549365	84147249	
Belgium	1951807	113031619.0	118195712	66372307	
Brazil	1504085	138276680.0	144697428	81387153	
Canada	3441944	306159372.0	319653077	180416413	
China	1842903	317244917.0	334082002	187417612	
Denmark	1602886	62013037.0	64526127	36803620	
Finland	1669339	188575323.0	198950287	111502578	
France	3211876	286569519.0	300166336	168780985	
Germany	3302882	262313078.0	275476051	154658565	
Italy	2514362	186648117.0	196181889	109579819	
Japan	3365197	357446635.0	374437913	210315931	
Korea	2074717	200725320.0	212939596	118797376	
Mexico	2165367	167187026.0	177346704	98849705	
Netherlands	2435565	184321687.0	193917160	108704507	
Singapore	2381155	197622402.0	208460742	116855753	
Spain	2312759	165066493.0	173636549	96906541	
Sweden	2093248	95411467.0	99623254	56122289	
Switzerland	1623527	100731878.0	104890676	58855881	
United Kingdom	2419810	277509546.0	291298415	163835204	
United States	3879316	826912620.0	866001599	487377157	

	Quantity	Unit_cost	Unit_price	Gross_profit	\
Retailer_country					
Australia	2000781	76472.0	139500	45045456.0	
Austria	2742824	95182.0	174512	59562183.0	
Belgium	2124791	75948.0	138325	46659336.0	
Brazil	2591989	65678.0	120706	56889537.0	
Canada	5722733	138890.0	252253	125742997.0	
China	6110945	89506.0	166923	129827307.0	
Denmark	1301136	58199.0	104951	25209421.0	
Finland	3603492	66496.0	120998	77072755.0	
France	5529613	135531.0	249099	117788568.0	
Germany	5084611	136106.0	249054	107654509.0	
Italy	3545695	104284.0	192412	77068308.0	
Japan	6787127	147334.0	273881	147130706.0	
Korea	3902092	90867.0	167671	81927974.0	
Mexico	3175752	91625.0	168281	68337351.0	
Netherlands	3448760	104459.0	192132	75617204.0	
Singapore	3788595	100989.0	186702	80766665.0	

Spain	3171715	100201.0	186009	68159942.0
Sweden	1681811	87679.0	159907	39289182.0
Switzerland	1822191	73297.0	137480	41876018.0
United Kingdom	5378361	101947.0	185999	113674387.0
United States	15722077	161142.0	294512	339535441.0

Retailer_country	Unit_sale_price	Revenue_missed_pct
Australia	132200.0	5.213563
Austria	164684.0	3.905030
Belgium	131558.0	4.369104
Brazil	114105.0	4.437362
Canada	239257.0	4.221359
China	157188.0	5.039806
Denmark	99282.0	3.894686
Finland	113849.0	5.214852
France	234453.0	4.529761
Germany	234939.0	4.778264
Italy	181504.0	4.859660
Japan	258109.0	4.537809
Korea	157702.0	5.736029
Mexico	158549.0	5.728710
Netherlands	180166.0	4.948233
Singapore	175607.0	5.199224
Spain	175661.0	4.935629
Sweden	151123.0	4.227715
Switzerland	130090.0	3.964888
United Kingdom	175439.0	4.733589
United States	278063.0	4.513731





- Which country had highest number of sales?
- Which product sold the most and. which product line had more revenue loss?
-

```
[297]: df.head(3)
```

```
[297]:   Year      Product_line Product_type      Product \
0  2015  Camping Equipment  Cooking Gear  TrailChef Water Bag
1  2015  Camping Equipment  Cooking Gear  TrailChef Water Bag
2  2015  Camping Equipment  Cooking Gear  TrailChef Water Bag

   Order_method_type Retailer_country  Revenue  Planned_revenue  Product_cost \
0      Telephone      United States  315044.0          437477      158372
1      Telephone           Canada    13445.0           14313        6299
2      Telephone           Japan   181120.0          235237      89413

   Quantity  Unit_cost  Unit_price  Gross_profit  Unit_sale_price
0     66385        3.0          7     156673.0          5.0
1      2172        3.0          7       7146.0          6.0
2     35696        3.0          7     91707.0          5.0
```

1.3 Feature selection

Feature selection huggely impacts performance of model. It : - reduces overfitting - improves accuracy - reduces training time

There are 3 methods of feature selection: - Univariate selection - Feature importance - Correlation matrix with Heatmap

```
[298]: # Adding more features

df['Missed_Revenue']=df['Planned_revenue']-df['Revenue']
df['revenue_to_cost_rate']=(df['Revenue']/df['Product_cost'])
```

```
[299]: df['revenue_to_cost_rate'].head(3)
```

```
[299]: 0    1.989266
      1    2.134466
      2    2.025656
      Name: revenue_to_cost_rate, dtype: float64
```

‘revenue_to_cost_rate’ of 1.989266 can be interpreted as - if you spend \$1, you can yield a revenue of \$1.989266. The more.. the better..

```
[300]: # Univariate selection

import numpy as np
from sklearn.feature_selection import SelectKBest
from sklearn.feature_selection import chi2

x = df.iloc[:,6:10] #independent columns
y = df.iloc[:,-2]   #target column i.e Gross_profit

#apply SelectKBest class to extract top 10 best features
bestfeatures = SelectKBest(score_func=chi2, k='all')
fit = bestfeatures.fit(x,y)
dfscores = pd.DataFrame(fit.scores_)
dfcolumns = pd.DataFrame(x.columns)


#concat two dataframes for better visualization
featureScores = pd.concat([dfcolumns,dfscores],axis=1)

featureScores.columns = ['Feature_Name','Score'] #naming the dataframe columns
print(featureScores.nlargest(10,'Score')) #print 10 best features
```

	Feature_Name	Score
1	Planned_revenue	1.131347e+10
0	Revenue	1.065375e+10
2	Product_cost	7.025373e+09
3	Quantity	2.488430e+08

```
[301]: # Let's see how Gross Profit is related with other variables using scatter plot.
import seaborn as sns
sns.pairplot(df, x_vars=['Product_cost', 'Planned_revenue', 'Revenue', 'Quantity'],
             y_vars=['Gross_profit', height=4, aspect=1, kind='scatter'])
```


```
plt.show()
```



Top three features that are selected by the algorithm does not have any outliers. However **Quantity** is not a linearly increasing metrics with **Gross_Profit**. We can ignore considering **Quantity** for the model, since we are taking the overall Revenue and Cost for determining the profits and not Unit_Cost or Unit_Sale

```
[302]: # Let's see the correlation between different variables.

fig, ax = plt.subplots(figsize=(10,10))
sns.heatmap(df.corr(), cmap="YlGnBu", annot = True, ax=ax)
plt.show()
```



This further proves that Revenue, Planned Revenue and Product Cost are the best and most correlated metrics to grossprofit

1.4 Predictive models

- Partition the data into 80% training and 20% validation.
- Explain how using validation set helps to avoid overfitting/underfitting
- At least build two models using decision tree, logistic regression, linear regression, neural network, and knn.
- Assess, analyze, and compare the performance of your models

1.4.1 Model 1: Multi-Variable Regression

```
[303]: X = df[['Revenue', 'Planned_revenue', 'Product_cost']]
      y = df['Gross_profit']
```

```
[304]: from sklearn.model_selection import train_test_split
      X_train, X_test, y_train, y_test = train_test_split(X, y, train_size = 0.8,
      ↪test_size = 0.2, random_state = 100)
```

```
[305]: X_train.head()
```

```
[305]:
```

	Revenue	Planned_revenue	Product_cost
13586	71571.0	73497	31013
14657	186462.0	196276	132296
2130	33813.0	35219	23594
23741	0.0	2840	1318
2486	13409.0	13822	8931

```
[306]: print(X_train.shape)
      print(X_test.shape)
      print(y_train.shape)
      print(y_test.shape)
```

```
(19794, 3)
(4949, 3)
(19794,)
(4949,)
```

```
[307]: import statsmodels.api as sm

      # Add a constant to get an intercept
      X_train_sm = sm.add_constant(X_train)

      # Fit the regression line using 'OLS'
      lr = sm.OLS(y_train, X_train_sm).fit()
```

```
[308]: lr.params
```

```
[308]: const          1.355763e-02
      Revenue        9.999998e-01
      Planned_revenue 6.130563e-08
      Product_cost    -9.999999e-01
      dtype: float64
```

```
[309]: # Performing a summary operation lists out all the different parameters of the
      ↪regression line fitted
      print(lr.summary())
```

OLS Regression Results

```

=====
Dep. Variable:          Gross_profit    R-squared:                1.000
Model:                  OLS            Adj. R-squared:          1.000
Method:                 Least Squares   F-statistic:              9.382e+14
Date:                   Tue, 10 Aug 2021 Prob (F-statistic):       0.00
Time:                   23:10:53        Log-Likelihood:           -10872.
No. Observations:       19794          AIC:                     2.175e+04
Df Residuals:           19790          BIC:                     2.178e+04
Df Model:                3
Covariance Type:        nonrobust
=====
===

```

	coef	std err	t	P> t	[0.025
0.975]					
const	0.0136	0.003	4.010	0.000	0.007
Revenue	1.0000	1.81e-07	5.51e+06	0.000	1.000
Planned_revenue	6.131e-08	1.7e-07	0.360	0.719	-2.73e-07
Product_cost	-1.0000	9.02e-08	-1.11e+07	0.000	-1.000

```

=====
Omnibus:                1271.077    Durbin-Watson:           1.995
Prob(Omnibus):           0.000      Jarque-Bera (JB):        5952.968
Skew:                    0.074      Prob(JB):                0.00
Kurtosis:                 5.683      Cond. No.:               7.67e+05
=====

```

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 7.67e+05. This might indicate that there are strong multicollinearity or other numerical problems.

Interpreting the model using OLS regression results:

- The coefficients and significance (p-values)
- R-squared
- F statistic and its significance

1. Coefficient Interpretation:

- Revenue The coefficient for Revenue is 1.0000, with a very low p value ($0.00 < 0.05$) The coefficient is positive, so it is positively correlated with Profit (Independent/ target variable) p value is statistically significant.

- **Planned_revenue** The coefficient for Planned_revenue is 6.131e-08, with a high p value (0.719 > 0.05) The coefficient is positive, so it is positively correlated with Profit (Independent/ target variable) p value is not statistically significant.
- **Product_cost** The coefficient for Product_cost is -1.0000, with a very low p value (0.00 < 0.05) It makes sense, as if there is increase in cost, the profits will decrease. SO it is negatively correlated. p value is statistically significant.

2. **R - squared is 1:**

- Meaning that 100% of the variance in Profit is explained by the 3 factors mentioned above
- This is a perfect R-squared value, which implies that independent variables and dependent variables are having a strong correlation.

3. **F statistic has a very low p value (practically low)**

- Meaning that the model fit is statistically significant, and the explained variance isn't purely by chance.

The fit is significant. Let us visualize how well the model fit the data. From our parameters, the linear regression equation is:

Gross_Profit = 0.0136 + 1.0000 * Revenue + 6.131e-08 * Planned_revenue - 1.0000 * Product_cost


```
[310]: y_train_pred = lr.predict(X_train_sm)
       res = (y_train - y_train_pred)
```

```
[311]: fig = plt.figure()
       sns.distplot(res, bins = 15)
       fig.suptitle('Error Terms', fontsize = 15)           # Plot heading
       plt.show()
```

/Users/Gaya/opt/anaconda3/lib/python3.8/site-packages/seaborn/distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

warnings.warn(msg, FutureWarning)

Error Terms



The residuals are following the normally distributed with a mean 0. All good!

```
[312]: X_test_sm = sm.add_constant(X_test)

# Predict the y values corresponding to X_test_sm
y_pred = lr.predict(X_test_sm)
```

```
[313]: y_pred.head()
```

```
[313]: 11845    2358.013430
      3353    3093.013437
      8863    77959.011071
      23037    4774.013349
      4364    2893.013698
      dtype: float64
```

```
[314]: from sklearn.metrics import mean_squared_error
      from sklearn.metrics import r2_score
```

```
[315]: #Returns the mean squared error; we'll take a square root
      np.sqrt(mean_squared_error(y_test, y_pred))
      r_squared = r2_score(y_test, y_pred)
      r_squared
```

```
[315]: 0.999999999992928
```

1.4.2 Model Summary

We can say that as there is an increase in the Cost, we are seeing good profits. Which means that whatever strategy that is being used for sales is clearly working good


1.5 Model 2: Logistic Regressionnn

```
[316]: df.revenue_to_cost_rate.describe()
```

```
[316]: count      24743.000000  
mean         1.908928  
std          0.497976  
min          0.000000  
25%         1.573688  
50%         1.829897  
75%         2.092941  
max          4.358086  
Name: revenue_to_cost_rate, dtype: float64
```

```
[317]: df['revenue_to_cost_rate'].plot.kde()
```

```
[317]: <AxesSubplot:ylabel='Density'>
```



Let us build a boolean value called 'profit_satisfied' with 'true' or 'false' values

Our threshold should be on the mean value here (1.9)

```
[318]: df['profit_satisfied'] = np.where(df['revenue_to_cost_rate']>1.9, 1, 0)
```

```
[319]: df.groupby('profit_satisfied')['profit_satisfied'].count()
```

```
[319]: profit_satisfied
0      13810
1       10933
Name: profit_satisfied, dtype: int64
```

Now the values are assigned, let us run a logistic regression to see if we can predict the sales satisfaction boolean We can create training and test dataset (80:20) as the previous model

```
[324]: X = df[['Revenue', 'Planned_revenue', 'Product_cost']]
y = df['profit_satisfied']
```

```
[325]: X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.
↪25,random_state=0)
```


```
[336]: from sklearn.linear_model import LogisticRegression
import seaborn as sn
from sklearn import metrics

logmodel = LogisticRegression()
logmodel.fit(X_train,y_train)

y_pred = logmodel.predict(X_test)
```

```
[337]: confusion_matrix = pd.crosstab(y_test, y_pred, rownames=['Actual'],
↪colnames=['Predicted'])
sn.heatmap(confusion_matrix, annot=True)
```

```
[337]: <AxesSubplot:xlabel='Predicted', ylabel='Actual'>
```



```
[338]: print('Accuracy: ',metrics.accuracy_score(y_test, y_pred))
plt.show()
```

Accuracy: 1.0

1.5.1 Model Summary

Accuracy of this logistic regression model is 100% meaning that **profit_satisfied** category completely depends on ['Revenue','Planned Revenue','Product Cost']

1.6 Future Steps & Conclusions

After detailed analysis we can say that the sales of outdoor products world wide is growing with the increase in the investments and some of the shortcomings of this analysis are

- Too many missing values
- More features on the data like customer related info, months, day, seasonality information
- Breakdown of cost and sales into multiple factors (operational cost, labour cost and so on)
- With more transaction-level customer metrics, we can build recommendation systems

If we can gather these details, we can expand our analysis to different dimensions